

## DEADLOCK: NECESSARY CONDITIONS

A deadlock is a situation where a set of processes is blocked because each process is holding a resource and waiting for another resource acquired by some other process. In this article, we will discuss deadlock, its necessary conditions, etc. in detail.

### What is Deadlock?

Deadlock is a situation in computing where two or more processes are unable to proceed because each is waiting for the other to release resources. Key concepts include mutual exclusion, resource holding, circular wait, and no preemption.


Consider an example when two trains are coming toward each other on the same track and there is only one track, none of the trains can move once they are in front of each other. This is a practical example of deadlock.

### How Does Deadlock occur in the Operating System?

Before going into detail about how deadlock occurs in the Operating System, let's first discuss how the Operating System uses the resources present. A process in an operating system uses resources in the following way.

- Requests a resource
- Use the resource
- Releases the resource

A situation occurs in operating systems when there are two or more processes that hold some resources and wait for resources held by other(s). For example, in the below diagram, Process 1 is holding Resource 1 and waiting for resource 2 which is acquired by process 2, and process 2 is waiting for resource 1.



### Examples of Deadlock

There are several examples of deadlock. Some of them are mentioned below.

1. The system has 2 tape drives. P0 and P1 each hold one tape drive and each needs another one.
2. Semaphores A and B, initialized to 1, P0, and P1 are in deadlock as follows:
  - P0 executes wait(A) and preempts.
  - P1 executes wait(B).
  - Now P0 and P1 enter in deadlock.

P0	P1
wait(A);	wait(B)
wait(B);	wait(A)

3. Assume the space is available for allocation of 200K bytes, and the following sequence of events occurs.

P0	P1
Request 80KB;	Request 70KB;
Request 60KB;	Request 80KB;

Deadlock occurs if both processes progress to their second request.

### **Necessary Conditions for Deadlock in OS**

Deadlock can arise if the following four conditions hold simultaneously (Necessary Conditions)

- **Mutual Exclusion:** Two or more resources are non-shareable (Only one process can use at a time).
- **Hold and Wait:** A process is holding at least one resource and waiting for resources.
- **No Preemption:** A resource cannot be taken from a process unless the process releases the resource.
- **Circular Wait:** A set of processes waiting for each other in circular form.

### **What are the Methods For Handling Deadlock?**

There are three ways to handle deadlock

- Deadlock Prevention or Avoidance
- Deadlock Recovery
- Deadlock Ignorance

<b>DEADLOCK PREVENTION METHODS</b>
------------------------------------

In deadlock prevention the aim is to not let full-fill one of the required condition of the deadlock. This can be done by this method:

#### **(i) Mutual Exclusion**

We only use the Lock for the non-share-able resources and if the resource is share-able (like read only file) then we not use the locks here. That ensure that in case of share -able resource, multiple process can access it at same time. Problem- Here the problem is that we can only do it in case of share-able resources but in case of no-share-able resources like printer , we have to use Mutual exclusion.

#### **(ii) Hold and Wait**

To ensure that Hold and wait never occurs in the system, we must guarantee that whenever process request for resource , it does not hold any other resources.

- We can provide the all resources to the process that is required for it's execution before starting it's execution . **problem** – for example if there are three resource that is required by a process and we have given all that resource before starting execution of process then there might be a situation that initially we required only two resource and after one hour we want third resources and this will cause starvation for the another process that wants this resources and in that waiting time that resource can allocated to other process and complete their execution.
- We can ensure that when a process request for any resources that time the process does not hold any other resources. Ex- Let there are three resources DVD, File and Printer . First the process request for DVD and File for the copying data into the file and let suppose it is going to take 1 hour and after it the process free all resources then again request for File and Printer to print that file.

### (iii) No Preemption

If a process is holding some resource and request other resources that are acquired and these resource are not available immediately then the resources that current process is holding are preempted. After some time process again request for the old resources and other required resources to re-start.

For example – Process p1 have resource r1 and requesting for r2 that is hold by process p2. then process p1 preempt r1 and after some time it try to restart by requesting both r1 and r2 resources.

Problem – This can cause the Live Lock Problem .

Live Lock : Live lock is the situation where two or more processes continuously changing their state in response to each other without making any real progress.

Example:

- suppose there are two processes p1 and p2 and two resources r1 and r2.
- Now, p1 acquired r1 and need r2 & p2 acquired r2 and need r1.

- so according to above method- Both p1 and p2 detect that they can't acquire second resource, so they release resource that they are holding and then try again.
- continuous cycle- p1 again acquired r1 and requesting to r2 p2 again acquired r2 and requesting to r1 so there is no overall progress still process are changing there state as they preempt resources and then again holding them. This the situation of Live Lock.

#### **(iv) Circular Wait:**

To remove the circular wait in system we can give the ordering of resources in which a process needs to acquire.

Ex: If there are process p1 and p2 and resources r1 and r2 then we can fix the resource acquiring order like the process first need to acquire resource r1 and then resource r2. so the process that acquired r1 will be allowed to acquire r2 , other process needs to wait until r1 is free.

This is the Deadlock prevention methods but practically only fourth method is used as all other three condition removal method have some disadvantages with them .

<b>DEADLOCK AVOIDANCE</b>
---------------------------

A deadlock avoidance policy grants a resource request only if it can establish that granting the request cannot lead to a deadlock either immediately or in the future. The kernal lacks detailed knowledge about future behavior of processes, so it cannot accurately predict deadlocks.


To facilitate deadlock avoidance under these conditions, it uses the following conservative approach: Each process declares the maximum number of resource units of each class that it may require. The kernal permits a process to request these resource units in stages- i.e. a few resource units at a time- subject to the maximum number declared by it and uses a worst case analysis technique to check for the possibility of future deadlocks. A request is granted only if there is no possibility of

deadlocks; otherwise, it remains pending until it can be granted. This approach is conservative because a process may complete its operation without requiring the maximum number of units declared by it.

Deadlock Avoidance is a process used by the Operating System to avoid Deadlock. Let's first understand what is Deadlock in an Operating System is. Deadlock is a situation that occurs in the Operating System when any Process enters a waiting state because another waiting process is holding the demanded resource. Deadlock is a common problem in multi-processing where several processes share a specific type of mutually exclusive resource known as a soft lock or software.

### How can an Operating System avoid Deadlock?

The operating system avoids Deadlock by knowing the maximum resource requirements of the processes initially, and also, the Operating System knows the free resources available at that time. The operating system tries to allocate the resources according to the process requirements and checks if the allocation can lead to a safe state or an unsafe state. If the resource allocation leads to an unsafe state, then the Operating System does not proceed further with the allocation sequence.



How does this happen? How does the Operating System allocate resources to the processes? We will try to understand the process of allocation of resources with an intuitive example further in this article.

## How does Deadlock Avoidance Work?

Let's understand the working of Deadlock Avoidance with the help of an intuitive example.

Process	Maximum Required	current Available	Need
P1	9	5	4
P2	5	2	3
P3	3	1	2

Let's consider three processes P1, P2, P3. Some more information on which the processes tell the Operating System are :

- P1 process needs a maximum of 9 resources (Resources can be any software or hardware Resources like tape drive or printer etc..) to complete its execution. P1 is currently allocated with 5 Resources and needs 4 more to complete its execution.
- P2 process needs a maximum of 5 resources and is currently allocated with 2 resources. So it needs 3 more resources to complete its execution.
- P3 process needs a maximum of 3 resources and is currently allocated with 1 resource. So it needs 2 more resources to complete its execution.
- The Operating System knows that only 2 resources out of the total available resources are currently free.

But only 2 resources are free now. Can P1, P2, and P3 satisfy their requirements? Let's try to find out.

As only 2 resources are free for now, only P3 can satisfy its need for 2 resources. If P3 takes 2 resources and completes its execution, then P3 can release its 3 (1+2) resources. Now the three free resources that P3 released can satisfy the need of P2. Now, P2 after taking the three free resources, can complete its execution and then release 5 (2+3) resources. Now five resources are free. P1 can now take 4 out of the 5 free resources and complete its execution. So, with 2 free resources available initially,

all the processes were able to complete their execution leading to a Safe State. The order of execution of the processes was  $\langle P3, P2, P1 \rangle$ .

What if initially there was only 1 free resource available? None of the processes would be able to complete its execution. Thus leading to an unsafe state.


We use two words, safe and unsafe states. What are those states? Let's understand these concepts.

### Safe State and Unsafe State

**Safe State** - In the above example, we saw that the Operating System was able to satisfy the needs of all three processes, P1, P2, and P3, with their resource requirements. So all the processes were able to complete their execution in a certain order like  $P3 \rightarrow P2 \rightarrow P1$ .

So, If the Operating System is able to allocate or satisfy the maximum resource requirements of all the processes in any order then the system is said to be in Safe State. So safe state does not lead to Deadlock.

**Unsafe State** - If the Operating System is not able to prevent Processes from requesting resources which can also lead to a Deadlock, then the System is said to be in an Unsafe State. Unsafe State does not necessarily cause deadlock it may or may not cause deadlock.



So, in the above diagram shows the three states of the System. An unsafe state does not always cause a Deadlock. Some unsafe states can lead to a Deadlock, as shown in the diagram.



### Deadlock Avoidance Example

Let's take an example that has multiple resource requirements for every Process. Let there be three Processes P1, P2, P3, and 4 resources R1, R2, R3, R4. The maximum resource requirements of the Processes are shown in the below table.

Process	R1	R2	R3	R4
P1	3	2	3	2
P2	2	3	1	4
P3	3	1	5	0

A number of currently allocated resources to the processes are:

Process	R1	R2	R3	R4
P1	1	2	3	1
P2	2	1	0	2
P3	2	0	1	0

The total number of resources in the System are :

R1	R2	R3	R4
7	4	5	4

We can find out the no of available resources for each of P1, P2, P3, P4 by subtracting the currently allocated resources from total resources.

Available Resources are :

R1	R2	R3	R4
2	1	1	1

Now, The need for the resources for the processes can be calculated by :

Need = Maximum Resources Requirement - Currently Allocated Resources.

The need for the Resources is shown below:

Process	R1	R2	R3	R4
P1	2	1	0	1
P2	0	2	1	2
P3	1	1	4	0

The available free resources are  $\langle 2,1,1,1 \rangle$  of resources of R1, R2, R3, and R4 respectively, which can be used to satisfy only the requirements of process P1 only initially as process P2 requires 2 R2 resources which are not available. The same is the case with Process P3, which requires 4 R3 resources which is not available initially.

The Steps for resources allotment is explained below:

1. Firstly, Process P1 will take the available resources and satisfy its resource need, complete its execution and then release all its allocated resources. Process P1 is initially allocated  $\langle 1,2,3,1 \rangle$  resources of R1, R2, R3, and R4 respectively. Process P1 needs  $\langle 2,1,0,1 \rangle$  resources of R1, R2, R3 and R4 respectively to complete its execution. So, process P1 takes the available free resources  $\langle 2,1,1,1 \rangle$  resources of R1, R2, R3, R4 respectively and can complete its execution and then release its current allocated resources and also the free resources it used to complete its execution. Thus P1 releases  $\langle 1+2,2+1,3+1,1+1 \rangle = \langle 3,3,4,2 \rangle$  resources of R1, R2, R3, and R4 respectively.
2. After step 1 now, available resources are now  $\langle 3,3,4,2 \rangle$ , which can satisfy the need of Process P2 as well as process P3. After process P2 uses the available Resources and completes its execution, the available resources are now  $\langle 5,4,4,4 \rangle$ .
3. Now, the available resources are  $\langle 5,4,4,4 \rangle$ , and the only Process left for execution is Process P3, which requires  $\langle 1,1,4,0 \rangle$  resources each of R1, R2, R3, and R4. So it can easily use the available resources and complete its execution. After P3 is executed, the resources available are  $\langle 7,4,5,4 \rangle$ , which is equal to the maximum resources or total resources available in the System.

So, the process execution sequence in the above example was <P1, P2, P3>. But it could also have been <P1, P3, P2> if process P3 would have been executed before process P2, which was possible as there were sufficient resources available to satisfy the need of both Process P2 and P3 after step 1 above.

## **BANKER'S ALGORITHM**

The Banker's Algorithm is a resource allocation and deadlock avoidance algorithm that tests for safety by simulating the allocation for the predetermined maximum possible amounts of all resources, then makes an "s-state" check to test for possible activities, before deciding whether allocation should be allowed to continue.

The Banker's Algorithm is a smart way for computer systems to manage how programs use resources, like memory or CPU time. It helps prevent situations where programs get stuck and can't finish their tasks, which is called deadlock. By keeping track of what resources each program needs and what's available, the algorithm makes sure that programs only get what they need in a safe order. This helps computers run smoothly and efficiently, especially when lots of programs are running at the same time.

### **Why Banker's Algorithm is Named So?**

The banker's algorithm is named so because it is used in the banking system to check whether a loan can be sanctioned to a person or not. Suppose there are  $n$  number of account holders in a bank and the total sum of their money is  $S$ . Let us assume that the bank has certain amount of money  $Y$ . If a person applies for a loan then the bank first subtracts the loan amount from the total money that the bank has ( $Y$ ) and if the remaining amount is greater than  $S$  then only the loan is sanctioned. It is done because if all the account holders come to withdraw their money then the bank can easily do it.

The Banker's Algorithm helps prevent deadlock by simulating resource allocation. If you want to understand this complex algorithm and its real-world applications, especially for GATE, the [GATE CS Self-Paced Course](#) provides a structured approach

to learning OS algorithms, including the Banker's Algorithm, with hands-on examples and problem-solving techniques

It also helps the OS to successfully share the resources between all the processes. It is called the banker's algorithm because bankers need a similar algorithm- they admit loans that collectively exceed the bank's funds and then release each borrower's loan in installments. The banker's algorithm uses the notation of a safe allocation state to ensure that granting a resource request cannot lead to a deadlock either immediately or in the future. In other words, the bank would never allocate its money in such a way that it can no longer satisfy the needs of all its customers. The bank would try to be in a safe state always.

The following Data structures are used to implement the Banker's Algorithm:

Let ' $n$ ' be the number of processes in the system and ' $m$ ' be the number of resource types.

#### **Available**

- It is a 1-d array of size ' $m$ ' indicating the number of available resources of each type.
- $Available[j] = k$  means there are ' $k$ ' instances of resource type  $R_j$

#### **Max**

- It is a 2-d array of size ' $n*m$ ' that defines the maximum demand of each process in a system.
- $Max[i, j] = k$  means process  $P_i$  may request at most ' $k$ ' instances of resource type  $R_j$ .

#### **Allocation**

- It is a 2-d array of size ' $n*m$ ' that defines the number of resources of each type currently allocated to each process.
- $Allocation[i, j] = k$  means process  $P_i$  is currently allocated ' $k$ ' instances of resource type  $R_j$

## Need

- It is a 2-d array of size ' $n \times m$ ' that indicates the remaining resource need of each process.
- $\text{Need}[i, j] = k$  means process  $P_i$  currently needs ' $k$ ' instances of resource type  $R_j$
- $\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$

Allocation specifies the resources currently allocated to process  $P_i$  and  $\text{Need}_i$  specifies the additional resources that process  $P_i$  may still request to complete its task.

Banker's algorithm consists of a Safety algorithm and a Resource request algorithm.

## Safety Algorithm

It is a safety algorithm used to check whether or not a system is in a safe state or follows the safe sequence in a banker's algorithm:

1. There are two vectors Work and Finish of length  $m$  and  $n$  in a safety algorithm.

Initialization

Initialize:  $\text{Work} = \text{Available}$

$\text{Finish}[i] = \text{false}$ ; for  $i = 0, 1, 2, 3, 4 \dots n - 1$ .

2. Check the availability status for each type of resources  $[i]$ , such as:

$\text{Need}[i] \leq \text{Work}$

$\text{Finish}[i] == \text{false}$

If the  $i$  does not exist, go to step 4.

3.  $\text{Work} = \text{Work} + \text{Allocation}(i)$  // to get new resource allocation

$\text{Finish}[i] = \text{true}$

Go to step 2 to check the status of resource availability for the next process.

4. If  $\text{Finish}[i] == \text{true}$ ; it means that the system is safe for all processes.

## Resource Request Algorithm

A resource request algorithm checks how a system will behave when a process makes each type of resource request in a system as a request matrix.

Let create a resource request array  $R[i]$  for each process  $P[i]$ . If the Resource Request  $i[j]$  equal to 'K', which means the process  $P[i]$  requires 'k' instances of Resources type  $R[j]$  in the system.

1. When the number of requested resources of each type is less than the Need resources, go to step 2 and if the condition fails, which means that the process  $P[i]$  exceeds its maximum claim for the resource. As the expression suggests:

If  $\text{Request}(i) \leq \text{Need}$

Go to step 2;

2. And when the number of requested resources of each type is less than the available resource for each process, go to step (3). As the expression suggests:

If  $\text{Request}(i) \leq \text{Available}$

Else Process  $P[i]$  must wait for the resource since it is not available for use.

3. When the requested resource is allocated to the process by changing state:

$\text{Available} = \text{Available} - \text{Request}$

$\text{Allocation}(i) = \text{Allocation}(i) + \text{Request}(i)$

$\text{Need}_i = \text{Need}_i - \text{Request}_i$

When the resource allocation state is safe, its resources are allocated to the process  $P(i)$ . And if the new state is unsafe, the Process  $P(i)$  has to wait for each type of Request  $R(i)$  and restore the old resource-allocation state.

### Example

Considering a system with five processes  $P_0$  through  $P_4$  and three resources of type A, B, C. Resource type A has 10 instances, B has 5 instances and type C has 7 instances. Suppose at time  $t_0$  following snapshot of the system has been taken:

Process	Allocation	Max	Available
	A B C	A B C	A B C
P <sub>0</sub>	0 1 0	7 5 3	3 3 2
P <sub>1</sub>	2 0 0	3 2 2	
P <sub>2</sub>	3 0 2	9 0 2	
P <sub>3</sub>	2 1 1	2 2 2	
P <sub>4</sub>	0 0 2	4 3 3	

### The content of the Need matrix

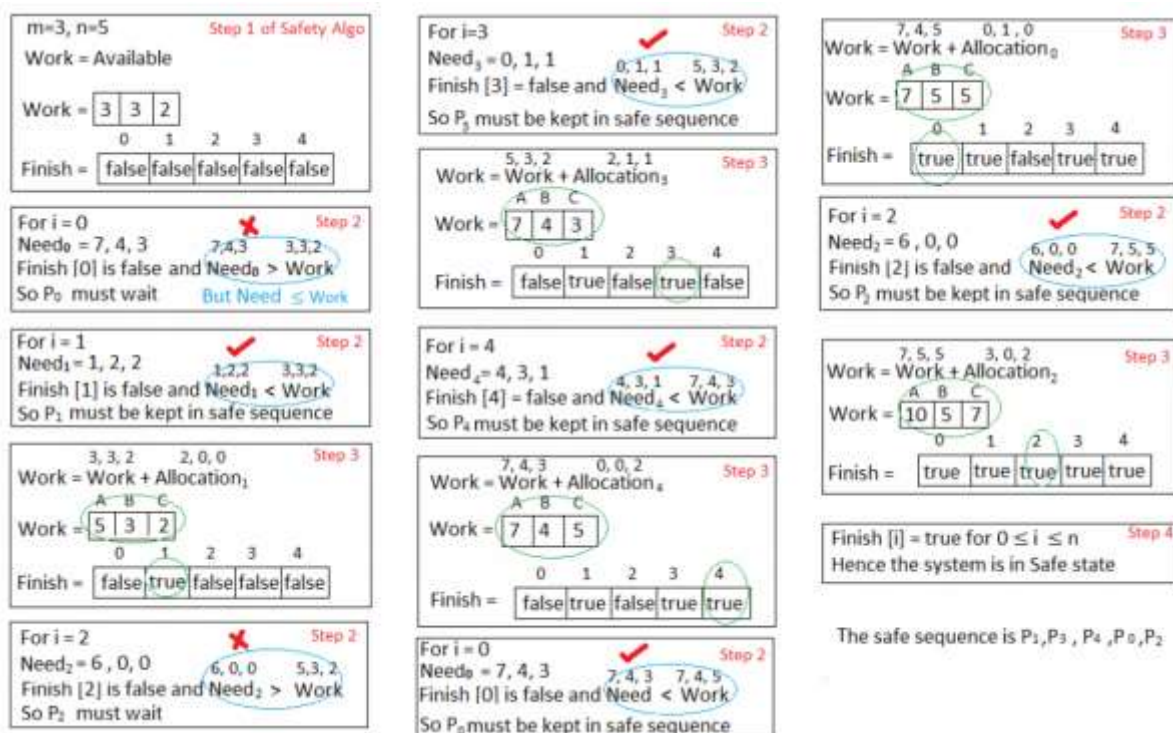
$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$

So, the content of Need Matrix is:

Process	Need		
	A	B	C
P <sub>0</sub>	7	4	3
P <sub>1</sub>	1	2	2
P <sub>2</sub>	6	0	0
P <sub>3</sub>	0	1	1
P <sub>4</sub>	4	3	1

We have to check whether the system is in a safe state or not? If Yes, then we have to find out the safe sequence


Applying the Safety algorithm on the given system,



What will happen if process P1 requests one additional instance of resource type A and two instances of resource type C?

Request<sub>1</sub> = 1, 0, 2

To decide whether the request is granted we use Resource Request algorithm



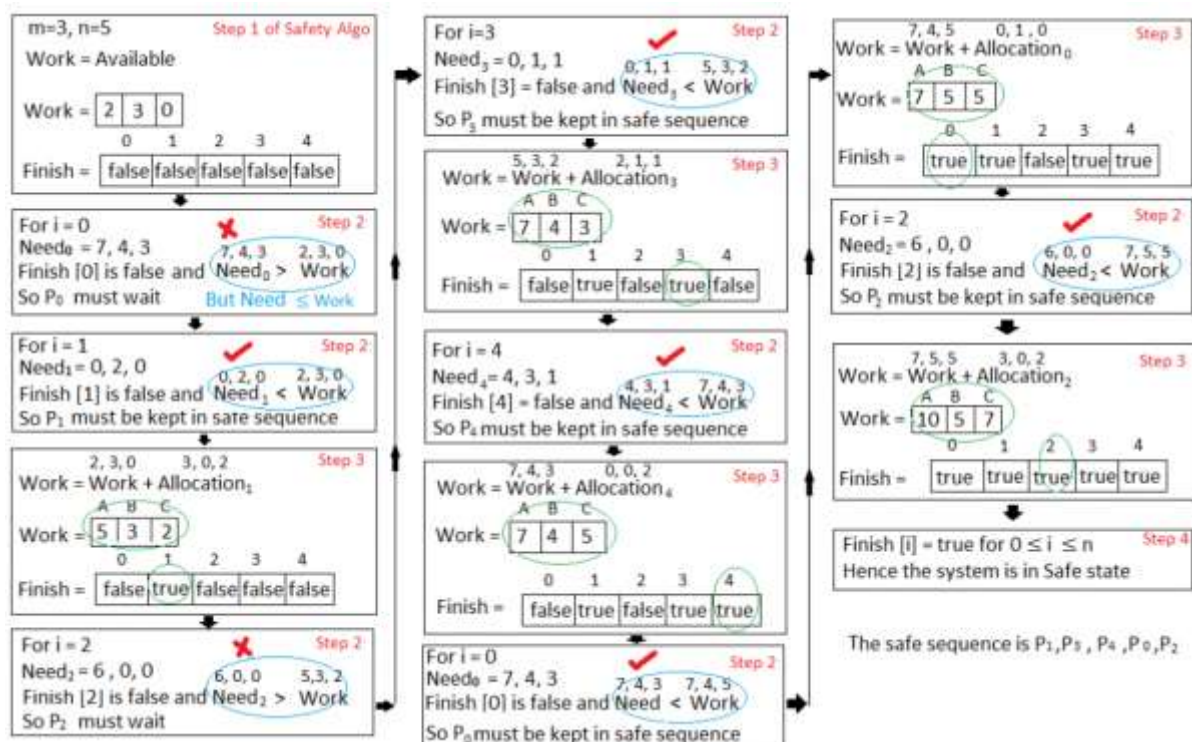
**Step 3**

Available = Available - Request<sub>1</sub>  
 Allocation<sub>1</sub> = Allocation<sub>1</sub> + Request<sub>1</sub>  
 Need<sub>1</sub> = Need<sub>1</sub> - Request<sub>1</sub>

Process	Allocation	Need	Available
	A B C	A B C	A B C
P <sub>0</sub>	0 1 0	7 4 3	2 3 0
P <sub>1</sub>	3 0 2	0 2 0	
P <sub>2</sub>	3 0 2	6 0 0	
P <sub>3</sub>	2 1 1	0 1 1	
P <sub>4</sub>	0 0 2	4 3 1	

We must determine whether this new system state is safe. To do so, we again execute Safety algorithm on the above data structures.





The Banker's Algorithm is a crucial method in operating system to ensure safe and efficient allocation of resources among processes. It helps prevent deadlock situations by carefully managing resource requests and releases. By keeping track of available resources and processes' needs, the algorithm ensures that resources are allocated in a way that avoids deadlock and maximizes system efficiency. Understanding and implementing the Banker's Algorithm is essential for maintaining system stability and ensuring that processes can complete their tasks without resource conflicts.

## DEADLOCK DETECTION AND RECOVERY

Deadlock Detection and Recovery is the mechanism of detecting and resolving deadlocks in an operating system. In operating systems, deadlock recovery is important to keep everything running smoothly. A deadlock occurs when two or more processes are blocked, waiting for each other to release the resources they need.

Detection methods help identify when this happens, and recovery techniques are used to resolve these issues and restore system functionality. This ensures that computers and devices can continue working without interruptions caused by deadlock situations. This can lead to a system-wide stall, where no process can make progress.

### **Approaches to Deadlock Detection and Recovery**

- **Prevention:** The operating system takes steps to prevent deadlocks from occurring by ensuring that the system is always in a safe state, where deadlocks cannot occur. This is achieved through resource allocation algorithms such as the Banker's Algorithm.
- **Detection and Recovery:** If deadlocks do occur, the operating system must detect and resolve them. Deadlock detection algorithms, such as the Wait-For Graph, are used to identify deadlocks, and recovery algorithms, such as the Rollback and Abort algorithm, are used to resolve them. The recovery algorithm releases the resources held by one or more processes, allowing the system to continue to make progress.

### **Difference Between Prevention and Detection/Recovery**


Prevention aims to avoid deadlocks altogether by carefully managing resource allocation, while detection and recovery aim to identify and resolve deadlocks that have already occurred.

Deadlock detection and recovery is an important aspect of operating system design and management, as it affects the stability and performance of the system. The choice of deadlock detection and recovery approach depends on the specific requirements of the system and the trade-offs between performance, complexity, and risk tolerance. The operating system must balance these factors to ensure that deadlocks are effectively detected and resolved.

## Deadlock Detection

### 1. If Resources Have a Single Instance

In this case for Deadlock detection, we can run an algorithm to check for the cycle in the Resource Allocation Graph. The presence of a cycle in the graph is a sufficient condition for deadlock.



In the above diagram, resource 1 and resource 2 have single instances. There is a cycle  $R1 \rightarrow P1 \rightarrow R2 \rightarrow P2$ . So, Deadlock is Confirmed.

### 2. If There are Multiple Instances of Resources

Detection of the cycle is necessary but not a sufficient condition for deadlock detection, in this case, the system may or may not be in deadlock varies according to different situations.

### 3. Wait-For Graph Algorithm

The Wait-For Graph Algorithm is a deadlock detection algorithm used to detect deadlocks in a system where resources can have multiple instances. The algorithm works by constructing a Wait-For Graph, which is a directed graph that represents the dependencies between processes and resources.

## Deadlock Recovery

A traditional operating system such as Windows doesn't deal with deadlock recovery as it is a time and space-consuming process. Real-time operating systems use Deadlock recovery.

- **Killing The Process:** Killing all the processes involved in the deadlock. Killing process one by one. After killing each process check for deadlock again and keep repeating the process till the system recovers from deadlock. Killing all the processes one by one helps a system to break circular wait conditions.
- **Resource Preemption:** Resources are preempted from the processes involved in the deadlock, and preempted resources are allocated to other processes so that there is a possibility of recovering the system from the deadlock. In this case, the system goes into starvation.
- **Concurrency Control:** Concurrency control mechanisms are used to prevent data inconsistencies in systems with multiple concurrent processes. These mechanisms ensure that concurrent processes do not access the same data at the same time, which can lead to inconsistencies and errors. Deadlocks can occur in concurrent systems when two or more processes are blocked, waiting for each other to release the resources they need. This can result in a system-wide stall, where no process can make progress. Concurrency control mechanisms can help prevent deadlocks by managing access to shared resources and ensuring that concurrent processes do not interfere with each other.

The Deadlock deduction algorithm employs several times varying data structures:

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to a process. The column represents resource and rows represent a process.

- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $request[i][j]$  equals  $k$  then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

Now, the Bankers algorithm includes a Safety Algorithm / Deadlock Detection Algorithm

The algorithm for finding out whether a system is in a safe state can be described as follows:

### Steps of Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$  respectively.  
Initialize *Work* = *Available*. For  $i=0, 1, \dots, n-1$ ,  
if *Request*  $i = 0$ , then *Finish* $[i] = \text{true}$ ; otherwise, *Finish* $[i] = \text{false}$ .
2. Find an index  $i$  such that both
  - a) *Finish* $[i] == \text{false}$
  - b) *Request*  $i \leq \text{Work}$
 If no such  $i$  exists go to step 4.
3. *Work* = *Work* + *Allocation*  $i$   
*Finish* $[i] = \text{true}$   
Go to Step 2.
4. If *Finish* $[i] == \text{false}$  for some  $i$ ,  $0 \leq i < n$ , then the system is in a deadlocked state.  
Moreover, if *Finish* $[i] == \text{false}$  the process  $P_i$  is deadlocked.

### Example:

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	0	0	0	0	0	0
P1	2	0	0	2	0	2			
P2	3	0	3	0	0	0			
P3	2	1	1	1	0	0			
P4	0	0	2	0	0	2			

- In this,  $Work = [0, 0, 0]$  &  
 $Finish = [false, false, false, false, false]$
- $i=0$  is selected as both  $Finish[0] = false$  and  $[0, 0, 0] \leq [0, 0, 0]$ .
- $Work = [0, 0, 0] + [0, 1, 0] \Rightarrow [0, 1, 0]$  &  
 $Finish = [true, false, false, false, false]$ .
- $i=2$  is selected as both  $Finish[2] = false$  and  $[0, 0, 0] \leq [0, 1, 0]$ .
- $Work = [0, 1, 0] + [3, 0, 3] \Rightarrow [3, 1, 3]$  &  
 $Finish = [true, false, true, false, false]$ .
- $i=1$  is selected as both  $Finish[1] = false$  and  $[2, 0, 2] \leq [3, 1, 3]$ .
- $Work = [3, 1, 3] + [2, 0, 0] \Rightarrow [5, 1, 3]$  &  
 $Finish = [true, true, true, false, false]$ .
- $i=3$  is selected as both  $Finish[3] = false$  and  $[1, 0, 0] \leq [5, 1, 3]$ .
- $Work = [5, 1, 3] + [2, 1, 1] \Rightarrow [7, 2, 4]$  &  
 $Finish = [true, true, true, true, false]$ .
- $i=4$  is selected as both  $Finish[4] = false$  and  $[0, 0, 2] \leq [7, 2, 4]$ .
- $Work = [7, 2, 4] + [0, 0, 2] \Rightarrow [7, 2, 6]$  &  
 $Finish = [true, true, true, true, true]$ .
- Since  $Finish$  is a vector of all true it means **there is no deadlock** in this example.

### Advantages of Deadlock Detection and Recovery

- **Improved System Stability:** Deadlocks can cause system-wide stalls, and detecting and resolving deadlocks can help to improve the stability of the system.
- **Better Resource Utilization:** By detecting and resolving deadlocks, the operating system can ensure that resources are efficiently utilized and that the system remains responsive to user requests.

- **Better System Design:** Deadlock detection and recovery algorithms can provide insight into the behavior of the system and the relationships between processes and resources, helping to inform and improve the design of the system.

### **Disadvantages of Deadlock Detection and Recovery**

- **Performance Overhead:** Deadlock detection and recovery algorithms can introduce a significant overhead in terms of performance, as the system must regularly check for deadlocks and take appropriate action to resolve them.
- **Complexity:** Deadlock detection and recovery algorithms can be complex to implement, especially if they use advanced techniques such as the Resource Allocation Graph or Timestamping.
- **False Positives and Negatives:** Deadlock detection algorithms are not perfect and may produce false positives or negatives, indicating the presence of deadlocks when they do not exist or failing to detect deadlocks that do exist.
- **Risk of Data Loss:** In some cases, recovery algorithms may require rolling back the state of one or more processes, leading to data loss or corruption.

Overall, the choice of deadlock detection and recovery approach depends on the specific requirements of the system, the trade-offs between performance, complexity, and accuracy, and the risk tolerance of the system. The operating system must balance these factors to ensure that deadlocks are effectively detected and resolved.

### **Conclusion**

This week we learned about the Deadlock and its Necessary Conditions, Deadlock Prevention Methods, Deadlock Avoidance, Bankers Algorithm with examples and Deadlock Detection and Recovery Techniques.